

## 7) Catálogo do Sistema

### 7.1) Tabelas de Sistema

#### 7.2) Informações sobre as tabelas de sistema

#### 7.3) Exemplos práticos

Todo SGBD precisa ter seu catálogo de sistema (metadados, dicionário de dados), onde armazena pelo menos:

- Nomes das tabelas
- Nomes dos campos
- Tipos de dados de cada campo
- As constraints
- Informações sobre os índices
- Privilégios de acesso dos elementos

O Catálogo de sistema está disponível no PostgreSQL desde a versão 7.4.

### 7.1) Tabelas de Sistema do PostgreSQL 8.3

Catalog Name	Purpose
<a href="#">pg_aggregate</a>	aggregate functions
<a href="#">pg_am</a>	index access methods
<a href="#">pg_amop</a>	access method operators
<a href="#">pg_amproc</a>	access method support procedures
<a href="#">pg_attrdef</a>	column default values
<a href="#">pg_attribute</a>	table columns ("attributes")
<a href="#">pg_authid</a>	authorization identifiers (roles)
<a href="#">pg_auth_members</a>	authorization identifier membership relationships
<a href="#">pg_autovacuum</a>	per-relation autovacuum configuration parameters
<a href="#">pg_cast</a>	casts (data type conversions)
<a href="#">pg_class</a>	tables, indexes, sequences, views ("relations")
<a href="#">pg_constraint</a>	check constraints, unique constraints, primary key constraints, foreign key constraints
<a href="#">pg_conversion</a>	encoding conversion information
<a href="#">pg_database</a>	databases within this database cluster
<a href="#">pg_depend</a>	dependencies between database objects
<a href="#">pg_description</a>	descriptions or comments on database objects

Catalog Name	Purpose
<a href="#">pg_enum</a>	enum label and value definitions
<a href="#">pg_index</a>	additional index information
<a href="#">pg_inherits</a>	table inheritance hierarchy
<a href="#">pg_language</a>	languages for writing functions
<a href="#">pg_largeobject</a>	large objects
<a href="#">pg_listener</a>	asynchronous notification support
<a href="#">pg_namespace</a>	schemas
<a href="#">pg_opclass</a>	access method operator classes
<a href="#">pg_operator</a>	operators
<a href="#">pg_opfamily</a>	access method operator families
<a href="#">pg_pltemplate</a>	template data for procedural languages
<a href="#">pg_proc</a>	functions and procedures
<a href="#">pg_rewrite</a>	query rewrite rules
<a href="#">pg_shdepend</a>	dependencies on shared objects
<a href="#">pg_shdescription</a>	comments on shared objects
<a href="#">pg_statistic</a>	planner statistics
<a href="#">pg_tablespace</a>	tablespaces within this database cluster
<a href="#">pg_trigger</a>	triggers
<a href="#">pg_ts_config</a>	text search configurations
<a href="#">pg_ts_config_map</a>	text search configurations' token mappings
<a href="#">pg_ts_dict</a>	text search dictionaries
<a href="#">pg_ts_parser</a>	text search parsers
<a href="#">pg_ts_template</a>	text search templates
<a href="#">pg_type</a>	data types

Mais detalhes em:

<http://www.postgresql.org/docs/8.3/interactive/catalogs.html>

<http://pgdocptbr.sourceforge.net/pg80/catalogs.html>

[http://www.cs.umu.se/kurser/TDBC86/H06/Slides/16\\_OH\\_catalog.pdf](http://www.cs.umu.se/kurser/TDBC86/H06/Slides/16_OH_catalog.pdf) (trazendo informações sobre o catálogo do PostgreSQL, Oracle e ODBC).

## 7.2) Informações sobre as tabelas de sistema

Para visualizar a estrutura de cada uma destas relações, podemos usar o psql. Exemplo:

```
\d pg_database
```

Observar que existe um esquema ()

**Para visualizar o catálogo completo, que está no esquema 'pg\_catalog', usar:**

```
\dS
```

Um total de 74 na versão 8.3 do PostgreSQL.

**Para visualizar a estrutura de uma das relações do esquema pg\_catalog:**

```
\d pg_catalog.pg_table
```

serão listados os campos, seus tipos e abaixo uma definição de view.

**Retornando todas as tabelas de sistema de um usuário:**

```
select * from pg_catalog.pg_tables where tableowner='dba1';
```

```
select * from pg_catalog.pg_tables where tableowner='postgres';
```

### 7.3) Exemplos práticos e úteis

#### Uso do disco pela Tabela

```
\c dba_projeto
```

```
vacuum analyze clientes;
```

```
SELECT relfilenode, relpages FROM pg_class WHERE relname = 'clientes';
```

relfilenode é o arquivo, com nome usando números num diretório do 'base'.

relpages é o número de páginas ocupado pela relação (tabela). Lembrar que cada página ocupa 8KB. relpages somente é atualizado por VACUUM, ANALYZE e uns poucos comandos de DDL como CREATE INDEX). O valor de relfilenode possui interesse caso se deseje examinar diretamente o arquivo em disco da tabela.

Podemos então saber algo com mais detalhes assim:

```
\c dba_projeto
```

```
vacuum analyze clientes;
```

```
SELECT relfilenode AS arquivo, relpages*8 AS tamanho_em_kb FROM pg_class WHERE  
relname = 'clientes';
```

No diretório do tutorial existe o arquivo [syscat.sql](http://pgdocptbr.sourceforge.net/pg80/syscat.sql) contendo várias consultas interessantes aos catálogos do sistema:

<http://pgdocptbr.sourceforge.net/pg80/syscat.sql>

```
-- syscat.sql-
```

```
-- Exemplos de consultas aos catálogos do sistema
```

```
-- Portions Copyright (c) 1996-2003, PostgreSQL Global Development Group
```

```
-- Portions Copyright (c) 1994, Regents of the University of California
```

```
-- Primeiro definir o caminho de procura do esquema como pg_catalog,
```

```
-- para não ser necessário qualificar todo objeto do sistema.
```

```
--
```

```
SET SEARCH_PATH TO pg_catalog;
```

```
--
```

```
-- Listar o nome de todos os administradores de banco de dados e o seus bancos de dados.
```

```
--
```

```
SELECT username, datname  
FROM pg_user, pg_database  
WHERE usesysid = datdba  
ORDER BY username, datname;
```

```
--
-- Listar todas as classes definidas pelo usuário
--
SELECT n.nspname, c.relname
FROM pg_class c, pg_namespace n
WHERE c.relnamespace=n.oid
      AND c.relkind = 'r'           -- sem índices, visões, etc
      AND n.nspname not like 'pg\\_%' -- sem catálogos
      AND n.nspname != 'information_schema' -- sem information_schema
ORDER BY nspname, relname;

-- Listar todos os índices simples (ou seja, àqueles que são definidos
-- sobre uma referência de coluna simples)
--
SELECT n.nspname AS schema_name,
       bc.relname AS class_name,
       ic.relname AS index_name,
       a.attname
FROM pg_namespace n,
     pg_class bc,      -- classe base
     pg_class ic,      -- classe índice
     pg_index i,
     pg_attribute a     -- atributo na base
WHERE bc.relnamespace = n.oid
      AND i.indrelid = bc.oid
      AND i.indexrelid = ic.oid
      AND i.indkey[0] = a.attnum
      AND i.indnatts = 1
      AND a.attrelid = bc.oid
ORDER BY schema_name, class_name, index_name, attname;

--
-- Listar os atributos definidos pelo usuário e seus tipos
-- para todas as classes definidas pelo usuário
--
SELECT n.nspname, c.relname, a.attname, format_type(t.oid, null) AS typname
FROM pg_namespace n, pg_class c,
     pg_attribute a, pg_type t
WHERE n.oid = c.relnamespace
      AND c.relkind = 'r'           -- sem índices
      AND n.nspname not like 'pg\\_%' -- sem catálogos
      AND n.nspname != 'information_schema' -- sem information_schema
      AND a.attnum > 0              -- sem atributo de sistema
      AND not a.attisdropped        -- sem colunas removidas
      AND a.attrelid = c.oid
      AND a.atttypid = t.oid
ORDER BY nspname, relname, attname;
```

```
--
-- Listar todos os tipos base definidos pelo usuário (sem incluir os tipos matriz)
--
```

```
SELECT n.nspname, u.username, format_type(t.oid, null) AS typename
FROM pg_type t, pg_user u, pg_namespace n
WHERE u.usesysid = t.typowner
      AND t.typnamespace = n.oid
      AND t.typrelid = '0'::oid      -- sem tipos complexos
      AND t.typelem = '0'::oid      -- sem matrizes
      AND n.nspname not like 'pg\\_%' -- sem catálogos
      AND n.nspname != 'information_schema' -- sem information_schema
ORDER BY nspname, username, typename;
```

```
--
-- Listar todas as funções de agregação e os tipos em que podem ser aplicadas
--
```

```
SELECT n.nspname, p.proname, format_type(t.oid, null) AS typename
FROM pg_namespace n, pg_aggregate a,
     pg_proc p, pg_type t
WHERE p.pronamespace = n.oid
      AND a.aggfnoid = p.oid
      AND p.proargtypes[0] = t.oid
ORDER BY nspname, proname, typename;
```

```
-- Restaurar o caminho de procura
--
```

```
RESET SEARCH_PATH;
```

### **Retornar o número de usuários conectados**

```
select count(*) from pg_stat_activity;
select count(*) from pg_stat_database;
```

pg\_stat\_database que apresenta para cada banco de dados o número de conexões.  
Fica mais fácil de visualizar do que o pg\_stat\_activity quando se tem muitas conexões.

### **Mostrar uso dos índices:**

```
select * from pg_statio_user_indexes;
select * from pg_stat_user_indexes;
```

### **Mostra estatística de uso de todas as tabelas e manutenção:**

```
select * from pg_stat_all_tables;
```

### **Mostra todas as tabelas e informações do atual esquema do atual banco:**

```
select * from pg_stat_user_tables;
```

Veja só o retorno: relid | schemaname | relname | seq\_scan | seq\_tup\_read | idx\_scan | idx\_tup\_fetch | n\_tup\_ins | n\_tup\_upd | n\_tup\_del | last\_vacuum | last\_autovacuum | last\_analyze | last\_autoanalyze

A função `pg_stat_get_backend_idset` provê uma conveniente maneira de gerar um registro/linha para cada processo ativo no servidor. Por exemplo, para **exibir os PIDs e as atuais consultas de todos os processos do servidor**:

```
SELECT pg_stat_get_backend_pid(s.backendid) AS procpid,  
       pg_stat_get_backend_activity(s.backendid) AS current_query  
FROM (SELECT pg_stat_get_backend_idset() AS backendid) AS s;
```

### Visualizar os processos do postgresql num UNIX:

```
ps auxww | grep ^post
```

### Mostrar todas as tabelas (inclusive de sistema) a quantidade de registros:

```
select relpages*8192 from pg_class;
```

## Funções para Administração do Sistema

Definindo configurações

A função `current_setting` retorna o valor corrente da definição `nome_da_definição`. Corresponde ao comando SQL `SHOW`. Por exemplo:

```
SELECT current_setting('datestyle');
```

A função `set_config` define o parâmetro `nome_da_configuração` como `novo_valor`. Se o parâmetro `é_local` for `true`, então o novo valor se aplica somente à transação corrente. Se for desejado que o novo valor seja aplicado à sessão corrente, deve ser utilizado `false`. Esta função corresponde ao comando SQL `SET`. Por exemplo:

```
SELECT set_config('log_statement_stats', 'off', false);
```

## Funções de Sinais para o Servidor

```
pg_cancel_backend(pid)  
pg_reload_conf()  
pg_rotate_logfile()
```

Se for bem-sucedida a função retorna 1, caso contrário retorna 0. O ID do processo (`pid`) de um servidor ativo pode ser encontrado a partir da coluna `procpid` da visão `pg_stat_activity`, ou listando os processos do postgres no servidor através do comando do Unix `ps`.

## Funções que Retornam o Tamanho de Objetos

Name	Return Type	Description
<code>pg_column_size(any)</code>	int	Number of bytes used to store a particular value (possibly compressed)
<code>pg_tablespace_size(oid)</code>	bigint	Disk space used by the tablespace with the specified OID
<code>pg_tablespace_size(name)</code>	bigint	Disk space used by the tablespace with the specified name
<code>pg_database_size(oid)</code>	bigint	Disk space used by the database with the specified OID
<code>pg_database_size(name)</code>	bigint	Disk space used by the database with the specified name
<code>pg_relation_size(oid)</code>	bigint	Disk space used by the table or index with the specified OID
<code>pg_relation_size(text)</code>	bigint	Disk space used by the table or index with the specified name. The table name may be qualified with a schema name
<code>pg_total_relation_size(oid)</code>	bigint	Total disk space used by the table with the specified OID, including indexes and toasted data
<code>pg_total_relation_size(text)</code>	bigint	Total disk space used by the table with the specified name, including indexes and toasted data. The table name may be qualified with a schema name
<code>pg_size_pretty(bi gint)</code>	text	Converts a size in bytes into a human-readable format with size units

`pg_column_size` shows the space used to store any individual data value.

`pg_tablespace_size` and `pg_database_size` accept the OID or name of a tablespace or database, and return the total disk space used therein.

`pg_relation_size` accepts the OID or name of a table, index or toast table, and returns the size in bytes.

`pg_total_relation_size` accepts the OID or name of a table or toast table, and returns the size in bytes of the data and all associated indexes and toast tables.

`pg_size_pretty` can be used to format the result of one of the other functions in a human-readable way, using kB, MB, GB or TB as appropriate.



## Funções de Acesso a Arquivos

Name	Return Type	Description
<code>pg_ls_dir(dirname text)</code>	setof text	List the contents of a directory
<code>pg_read_file(filename text, offset bigint, length bigint)</code>	text	Return the contents of a text file
<code>pg_stat_file(filename text)</code>	record	Return information about a file

`pg_ls_dir` returns all the names in the specified directory, except the special entries "." and "..".

`pg_read_file` returns part of a text file, starting at the given `offset`, returning at most `length` bytes (less if the end of file is reached first). If `offset` is negative, it is relative to the end of the file.

`pg_stat_file` returns a record containing the file size, last accessed time stamp, last modified time stamp, last file status change time stamp (Unix platforms only), file creation timestamp (Windows only), and a boolean indicating if it is a directory. Typical usages include:

```
SELECT * FROM pg_stat_file('filename');
SELECT (pg_stat_file('filename')).modification;
```

Mais detalhes em: <http://www.postgresql.org/docs/8.3/interactive/functions-admin.html>

## Consultando a Estrutura de uma Tabela através do Catálogo

SELECT

rel.nspname, rel.relname, attrs.attname, "Type", "Default", attrs.attnotnull

FROM (

SELECT c.oid, n.nspname, c.relname

FROM pg\_catalog.pg\_class c

LEFT JOIN pg\_catalog.pg\_namespace n ON n.oid = c.relnamespace

WHERE pg\_catalog.pg\_table\_is\_visible(c.oid) ) rel

JOIN (

SELECT a.attname, a.attrelid, pg\_catalog.format\_type(a.atttypid, a.atttypmod)

as "Type",

(SELECT substring(d.adsrc for 128) FROM pg\_catalog.pg\_attrdef d

WHERE d.adrelid = a.attrelid AND d.adnum = a.attnum AND a.atthasdef)

as "Default",

a.attnotnull, a.attnum

FROM pg\_catalog.pg\_attribute a WHERE a.attnum > 0 AND NOT a.attisdropped )

attrs

ON (attrs.attrelid = rel.oid )

WHERE relname = 'clientes' ORDER BY attrs.attnum;

Função em PLPGSQL:

```
CREATE OR REPLACE FUNCTION Dados_Tabela(varchar(30))
RETURNS SETOF tabela_estrutura AS '
DECLARE
  r tabela_estrutura%ROWTYPE;
  rec RECORD;
  vTabela alias for $1;
  eSql TEXT;

BEGIN
  eSql := 'SELECT
    CAST(rel.nspname as TEXT), CAST(rel.relname AS TEXT) ,
    CAST(attrs.attname AS TEXT), CAST("Type" AS TEXT),
    CAST("Default" AS TEXT), attrs.attnotnull
  FROM
    (SELECT c.oid, n.nspname, c.relname
     FROM pg_catalog.pg_class c
     LEFT JOIN pg_catalog.pg_namespace n ON n.oid = c.relnamespace
     WHERE pg_catalog.pg_table_is_visible(c.oid) ) rel
  JOIN
    (SELECT a.attname, a.attrelid,
     pg_catalog.format_type(a.atttypid, a.atttypmod) as "Type",
     (SELECT substring(d.adsrc for 128) FROM
pg_catalog.pg_attrdef d
     WHERE d.adrelid = a.attrelid AND d.adnum = a.attnum AND
a.atthasdef)
    as "Default", a.attnotnull, a.attnum
    FROM pg_catalog.pg_attribute a
    WHERE a.attnum > 0 AND NOT a.attisdropped ) attrs
    ON (attrs.attrelid = rel.oid )
    WHERE relname LIKE '||vTabela||' || '%' || vTabela || '%' ||
    ORDER BY attrs.attnum';
  FOR r IN EXECUTE eSql
  LOOP
    RETURN NEXT r;
  END LOOP;
  IF NOT FOUND THEN
    RAISE EXCEPTION 'Tabela % não encontrada', vTabela;
  END IF;
  RETURN;
END
';
LANGUAGE 'plpgsql';
```

Usando:

```
SELECT * FROM Dados_Tabela('clientes');
```

Fonte: [http://imasters.uol.com.br/artigo/2137/postgresql/consultando\\_a\\_estrutura\\_de\\_uma\\_tabela/](http://imasters.uol.com.br/artigo/2137/postgresql/consultando_a_estrutura_de_uma_tabela/)

### **Retornando Informações sobre uma Tabela**

```
SELECT pg_attribute.attnum AS index,  
attname AS field,  
typename AS type,  
atttypmod-4 as length,  
NOT attnotnull AS "null",  
adsrc AS def  
FROM pg_attribute,  
pg_class,  
pg_type,  
pg_attrdef  
WHERE pg_class.oid=attrelid  
AND pg_type.oid=atttypid  
AND attnum>0  
AND pg_class.oid=adrelid  
AND adnum=attnum  
AND atthasdef='t'  
AND lower(relname)='clientes'  
UNION  
SELECT pg_attribute.attnum AS index,  
attname AS field,  
typename AS type,  
atttypmod-4 as length,  
NOT attnotnull AS "null",  
" AS def  
FROM pg_attribute,  
pg_class,  
pg_type  
WHERE pg_class.oid=attrelid  
AND pg_type.oid=atttypid
```

```
AND attnum>0
AND atthasdef='f'
AND lower(relname)='clientes';
```

Use o comando \x no psql antes de executar a consulta caso queira exibir os registros em sequência, \x novamente para voltar ao normal.

Fonte:

[http://imasters.uol.com.br/artigo/1283/postgresql/informacoes\\_atraves\\_do\\_catalogo\\_do\\_sistema/](http://imasters.uol.com.br/artigo/1283/postgresql/informacoes_atraves_do_catalogo_do_sistema/)

Podemos listar tabelas, índices, sequências e vies usando os comandos do psql:

```
\d{t|i|s|v}
```

Abaixo veremos mais algumas funções que, ao contrário, usarão o SQL para receber essas informações.

Execute o trecho do script dba\_projeto para criação do banco 'catalogo' da Aula 8.

### Listando Tabelas do Banco Atual

```
SELECT relname
  FROM pg_class
 WHERE relname !~ '^(pg_|sql_) '
    AND relkind = 'r';

-- using INFORMATION_SCHEMA:

SELECT table_name
  FROM information_schema.tables
 WHERE table_type = 'BASE TABLE'
    AND table_schema NOT IN
      ('pg_catalog', 'information_schema');
```

### Listando Views do Banco Atual

```
-- with postgresql 7.2:

SELECT viewname
  FROM pg_views
 WHERE viewname !~ '^pg_';

-- with postgresql 7.4 and later:

SELECT viewname
  FROM pg_views
 WHERE schemaname NOT IN
      ('pg_catalog', 'information_schema')
    AND viewname !~ '^pg_';

-- using INFORMATION_SCHEMA:
```

```
SELECT table_name
FROM information_schema.tables
WHERE table_type = 'VIEW'
AND table_schema NOT IN
    ('pg_catalog', 'information_schema')
AND table_name !~ '^pg_';

-- or

SELECT table_name
FROM information_schema.views
WHERE table_schema NOT IN ('pg_catalog', 'information_schema')
AND table_name !~ '^pg_';
```

### Listando Todos os Usuários

```
SELECT username
FROM pg_user;
```

### Retornando os nomes dos Campos de uma Tabela

```
SELECT a.attname
FROM pg_class c, pg_attribute a, pg_type t
WHERE c.relname = 'test2'
AND a.attnum > 0
AND a.attrelid = c.oid
AND a.atttypid = t.oid;
```

-- Usando INFORMATION\_SCHEMA:

```
SELECT column_name
FROM information_schema.columns
WHERE table_name = 'test2';
```

### Informações Detalhadas sobre os Campos de uma Tabela

```
SELECT a.attnum AS ordinal_position,
       a.attname AS column_name,
       t.typname AS data_type,
       a.attlen AS character_maximum_length,
       a.atttypmod AS modifier,
       a.attnotnull AS notnull,
       a.atthasdef AS hasdefault,
       col_description(a.attrelid, a.attnum) as field_comment
FROM pg_class c,
     pg_attribute a,
     pg_type t
WHERE c.relname = 'test2'
AND a.attnum > 0
AND a.attrelid = c.oid
AND a.atttypid = t.oid
ORDER BY a.attnum;
```

-- Usando INFORMATION\_SCHEMA:

```
SELECT ordinal_position,
       column_name,
       data_type,
       column_default,
       is_nullable,
       character_maximum_length,
       numeric_precision
FROM information_schema.columns
WHERE table_name = 'test2'
ORDER BY ordinal_position;
```

### Retornando os Nomes de Índices de uma Tabela

```
SELECT relname
FROM pg_class
WHERE oid IN (
    SELECT indexrelid
    FROM pg_index, pg_class
    WHERE pg_class.relname='test2'
    AND pg_class.oid=pg_index.indrelid
    AND indisunique != 't'
    AND indisprimary != 't'
);
```

### Retornando as Constraints de uma Tabela

```
SELECT conname
FROM pg_constraint, pg_class
WHERE pg_constraint.conrelid = pg_class.oid
AND relname = 'test2';
```

-- with INFORMATION\_SCHEMA:

```
SELECT constraint_name, constraint_type
FROM information_schema.table_constraints
WHERE table_name = 'test2';
```

### Recebendo Informações Detalhadas sobre as Constraints de uma Tabela

```
SELECT c.conname AS constraint_name,
       CASE c.contype
         WHEN 'c' THEN 'CHECK'
         WHEN 'f' THEN 'FOREIGN KEY'
         WHEN 'p' THEN 'PRIMARY KEY'
         WHEN 'u' THEN 'UNIQUE'
       END AS "constraint_type",
       CASE WHEN c.condeferrable = 'f' THEN 0 ELSE 1 END AS is_deferrable,
       CASE WHEN c.condeferred = 'f' THEN 0 ELSE 1 END AS is_deferred,
       t.relname AS table_name,
       array_to_string(c.conkey, ' ') AS constraint_key,
       CASE confupdtype
         WHEN 'a' THEN 'NO ACTION'
         WHEN 'r' THEN 'RESTRICT'
         WHEN 'c' THEN 'CASCADE'
       END AS confupdtype
```

```

        WHEN 'n' THEN 'SET NULL'
        WHEN 'd' THEN 'SET DEFAULT'
    END AS on_update,
    CASE confdeltype
        WHEN 'a' THEN 'NO ACTION'
        WHEN 'r' THEN 'RESTRICT'
        WHEN 'c' THEN 'CASCADE'
        WHEN 'n' THEN 'SET NULL'
        WHEN 'd' THEN 'SET DEFAULT'
    END AS on_delete,
    CASE confmatchtype
        WHEN 'u' THEN 'UNSPECIFIED'
        WHEN 'f' THEN 'FULL'
        WHEN 'p' THEN 'PARTIAL'
    END AS match_type,
    t2.relname AS references_table,
    array_to_string(c.confkey, ' ') AS fk_constraint_key
FROM pg_constraint c
LEFT JOIN pg_class t ON c.conrelid = t.oid
LEFT JOIN pg_class t2 ON c.confrelid = t2.oid
WHERE t.relname = 'testconstraints2'
AND c.conname = 'testconstraints_id_fk';

```

```
-- with INFORMATION_SCHEMA:
```

```

SELECT tc.constraint_name,
       tc.constraint_type,
       tc.table_name,
       kcu.column_name,
       tc.is_deferrable,
       tc.initially_deferred,
       rc.match_option AS match_type,
       rc.update_rule AS on_update,
       rc.delete_rule AS on_delete,
       ccu.table_name AS references_table,
       ccu.column_name AS references_field
FROM information_schema.table_constraints tc
LEFT JOIN information_schema.key_column_usage kcu
    ON tc.constraint_catalog = kcu.constraint_catalog
    AND tc.constraint_schema = kcu.constraint_schema
    AND tc.constraint_name = kcu.constraint_name
LEFT JOIN information_schema.referential_constraints rc
    ON tc.constraint_catalog = rc.constraint_catalog
    AND tc.constraint_schema = rc.constraint_schema
    AND tc.constraint_name = rc.constraint_name
LEFT JOIN information_schema.constraint_column_usage ccu
    ON rc.unique_constraint_catalog = ccu.constraint_catalog
    AND rc.unique_constraint_schema = ccu.constraint_schema
    AND rc.unique_constraint_name = ccu.constraint_name
WHERE tc.table_name = 'testconstraints2'
AND tc.constraint_name = 'testconstraints_id_fk';

```

## Listando as Sequências

```

SELECT relname
FROM pg_class
WHERE relkind = 'S'
AND relnamespace IN (

```

```
SELECT oid
FROM pg_namespace
WHERE nspname NOT LIKE 'pg_%'
AND nspname != 'information_schema'
);
```

### Listando Todas as Triggers

```
SELECT trg.tgname AS trigger_name
FROM pg_trigger trg, pg_class tbl
WHERE trg.tgrelid = tbl.oid
AND tbl.relname !~ '^pg_';
-- or
SELECT tgname AS trigger_name
FROM pg_trigger
WHERE tgname !~ '^pg_';

-- with INFORMATION_SCHEMA:

SELECT DISTINCT trigger_name
FROM information_schema.triggers
WHERE trigger_schema NOT IN
('pg_catalog', 'information_schema');
```

### Listando Somente as Triggers de uma Tabela

```
SELECT trg.tgname AS trigger_name
FROM pg_trigger trg, pg_class tbl
WHERE trg.tgrelid = tbl.oid
AND tbl.relname = 'newtable';

-- with INFORMATION_SCHEMA:

SELECT DISTINCT trigger_name
FROM information_schema.triggers
WHERE event_object_table = 'newtable'
AND trigger_schema NOT IN
('pg_catalog', 'information_schema');
```

### Recebendo Informações Detalhadas sobre as Triggers

```
SELECT trg.tgname AS trigger_name,
tbl.relname AS table_name,
p.proname AS function_name,
CASE trg.tgtype & cast(2 as int2)
WHEN 0 THEN 'AFTER'
ELSE 'BEFORE'
END AS trigger_type,
CASE trg.tgtype & cast(28 as int2)
WHEN 16 THEN 'UPDATE'
WHEN 8 THEN 'DELETE'
WHEN 4 THEN 'INSERT'
WHEN 20 THEN 'INSERT, UPDATE'
```



```

        WHEN 28 THEN 'INSERT, UPDATE, DELETE'
        WHEN 24 THEN 'UPDATE, DELETE'
        WHEN 12 THEN 'INSERT, DELETE'
    END AS trigger_event
FROM pg_trigger trg,
     pg_class tbl,
     pg_proc p
WHERE trg.tgrelid = tbl.oid
     AND trg.tgfoid = p.oid
     AND tbl.relname !~ '^pg_';

-- with INFORMATION_SCHEMA:

SELECT *
FROM information_schema.triggers
WHERE trigger_schema NOT IN
      ('pg_catalog', 'information_schema');
```

## Listar as Funções

```

SELECT proname
FROM pg_proc pr,
     pg_type tp
WHERE tp.oid = pr.prorettype
     AND pr.proisagg = FALSE
     AND tp.typname <> 'trigger'
     AND pr.pronamespace IN (
        SELECT oid
        FROM pg_namespace
        WHERE nspname NOT LIKE 'pg_%'
          AND nspname != 'information_schema'
    );

-- with INFORMATION_SCHEMA:

SELECT routine_name
FROM information_schema.routines
WHERE specific_schema NOT IN
      ('pg_catalog', 'information_schema')
     AND type_udt_name != 'trigger';
```

## Outra mais detalhada:

```

CREATE OR REPLACE FUNCTION public.function_args(
    IN funcname character varying,
    IN schema character varying,
    OUT pos integer,
    OUT direction character,
    OUT argname character varying,
    OUT datatype character varying)
RETURNS SETOF RECORD AS $$DECLARE
    rettype character varying;
    argtypes oidvector;
    allargtypes oid[];
    argmodes "char"[];
```

```

argnames text[];
mini integer;
maxi integer;
BEGIN
/* get object ID of function */
SELECT INTO rettype, argtypes, allargtypes, argmodes, argnames
CASE
WHEN pg_proc.proretset
THEN 'setof ' || pg_catalog.format_type(pg_proc.prorettype, NULL)
ELSE pg_catalog.format_type(pg_proc.prorettype, NULL) END,
pg_proc.proargtypes,
pg_proc.proallargtypes,
pg_proc.proargmodes,
pg_proc.proargnames
FROM pg_catalog.pg_proc
JOIN pg_catalog.pg_namespace
ON (pg_proc.pronamespace = pg_namespace.oid)
WHERE pg_proc.prorettype <> 'pg_catalog.cstring'::pg_catalog.regtype
AND (pg_proc.proargtypes[0] IS NULL
OR pg_proc.proargtypes[0] <> 'pg_catalog.cstring'::pg_catalog.regtype)
AND NOT pg_proc.proisagg
AND pg_proc.proname = funcname
AND pg_namespace.nspname = schema
AND pg_catalog.pg_function_is_visible(pg_proc.oid);

/* bail out if not found */
IF NOT FOUND THEN
RETURN;
END IF;

/* return a row for the return value */
pos = 0;
direction = 'o'::char;
argname = 'RETURN VALUE';
datatype = rettype;
RETURN NEXT;

/* unfortunately allargtypes is NULL if there are no OUT parameters */
IF allargtypes IS NULL THEN
mini = array_lower(argtypes, 1); maxi = array_upper(argtypes, 1);
ELSE
mini = array_lower(allargtypes, 1); maxi = array_upper(allargtypes, 1);
END IF;
IF maxi < mini THEN RETURN; END IF;

/* loop all the arguments */
FOR i IN mini .. maxi LOOP
pos = i - mini + 1;
IF argnames IS NULL THEN
argname = NULL;
ELSE
argname = argnames[i];
END IF;
IF allargtypes IS NULL THEN
direction = 'i'::char;
datatype = pg_catalog.format_type(argtypes[i], NULL);
ELSE
direction = argmodes[i];
datatype = pg_catalog.format_type(allargtypes[i], NULL);

```

```

END IF;
RETURN NEXT;
END LOOP;

RETURN;
END;$$ LANGUAGE plpgsql STABLE STRICT SECURITY INVOKER;
COMMENT ON FUNCTION public.function_args(character varying, character
varying)
IS $$For a function name and schema, this procedure selects for each
argument the following data:
- position in the argument list (0 for the return value)
- direction 'i', 'o', or 'b'
- name (NULL if not defined)
- data type$$;

```

Do artigo de Lorenzo Albeton. Extracting metadata from PostgreSQL using INFORMATION\_SCHEMA :

[http://www.alberton.info/postgresql\\_meta\\_info.html](http://www.alberton.info/postgresql_meta_info.html)

Que também traz artigos sobre os catálogos do Firebird, Oracle e SQL Server.

### **Saber quantidade de registros no banco inteiro:**

```
SELECT sum(C.reltuples)::int FROM pg_class C WHERE c.relkind = 'r':"char";
```

Fonte: Blog da Kenia Milene

<http://keniamilene.wordpress.com/2007/09/18/contar-registros-em-banco-postgresql/>

### **Listando os bancos e sua codificação**

```
SELECT datname, pg_encoding_to_char(encoding) FROM pg_database;
```

### **Exibindo codificações e versões**

```
SELECT name, setting FROM pg_settings
WHERE name ~ 'encoding|^lc_|version';
```

### **Mostrar data da última execução do Vacuum e do Autovacuum**

```
select relname, last_vacuum, last_autovacuum from
pg_stat_all_tables where schemaname like 'public';
```

Vide o capítulo "Table Statistics" do Livro PostgreSQL The Comprehensive Guide, da Sam's em:

<http://www.iphelp.ru/faq/15/ch04lev1sec4.html>

**Exibir Todas as Tabelas e seus Registros**

```
select
    n.nspname as esquema,
    c.relname as tabela, c.reltuples::int as registros
from pg_class c
    left join pg_namespace n on n.oid = c.relnamespace
    left join pg_tablespace t on t.oid = c.reltablespace
where c.relkind = 'r'::char
    and nspname not in('information_schema','pg_catalog', 'pg_toast')
order by n.nspname,registros;
```

Outra:

```
select relname as tabelas, reltuples
from pg_class
where relkind = 'r'::char and relname not like 'pg_%' and relname not like 'sql_%' order by reltuples;
```

Outra solução:

VACUUM ANALYZE (em todo o banco)

```
SELECT sum(reltuples) as qtd_linhas
FROM pg_class
WHERE relkind = 'r';
```

criar uma view com a instrução SQL

Rodrigo Hjort - <http://icewall.org/~hjort> – na lista pgbr-geral

Boa fonte de leitura sobre o Catálogo do Sistema do PostgreSQL:

System Tables no capítulo 6 do livro: PostgreSQL Developer's Handbook de Ewald Geschwinde, Hans-Jürgen Schönig.